

# Cluster Computing

Dan Barker [danbarke@umich.edu](mailto:danbarke@umich.edu)  
Cluster System Administrator

# Batch Computing Clusters

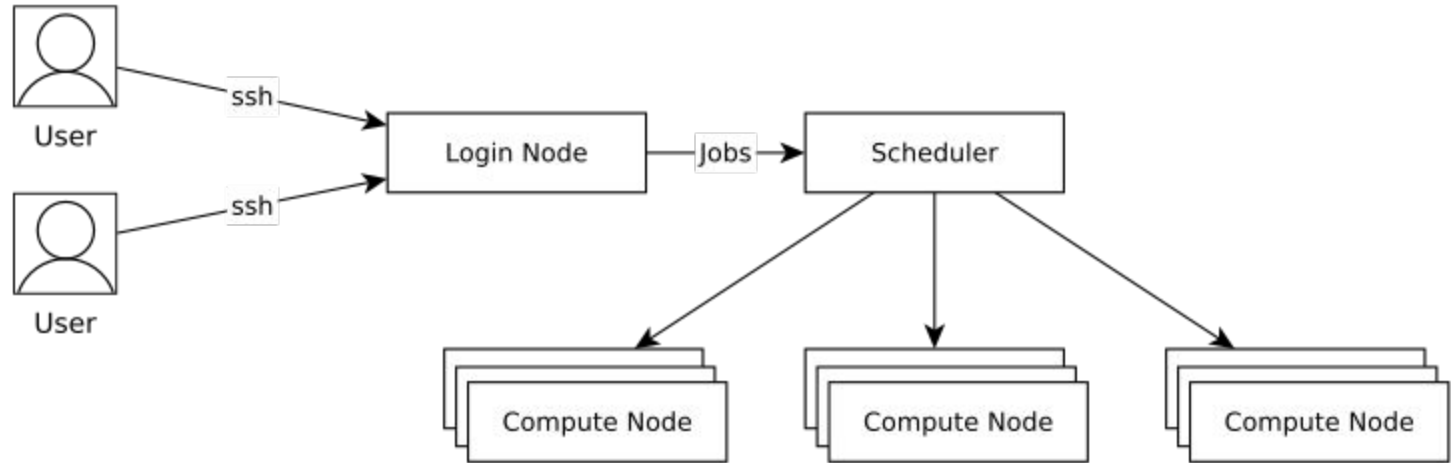


image credit: <https://docs.hpc.qmul.ac.uk/intro/>

# Biostatistics Computing Cluster

55 Compute Nodes

2 Login Nodes

948 Processors

~5.5TB RAM

Red Hat Enterprise Linux

Statistical Software R, SAS, Matlab

# Login with SSH

login host: `biostat-login.sph.umich.edu`

Enter kerberos [password](#) (same as email)

Authenticate with [Duo](#)

MacOS: [Terminal](#)

Windows: [PuTTY](#)

All OSes: [Biostatistics Web Portal](#)

# Anatomy of a Batch Job

```
#!/bin/bash
#SBATCH --job-name=hello_world
#SBATCH --time=10:00
#SBATCH --mail-user=danbarke@umich.edu
#SBATCH --mail-type=BEGIN,END,FAIL
#SBATCH --mem=4000mb
#SBATCH --cpus-per-task=1
```

```
R CMD BATCH script.R
```

# Slurm Commands

`sbatch <batch_script_file>` - submits a job

`squeue -u <username>` - show all jobs for user on cluster

`scancel <jobid>` - cancel a job

`sacct` - show data about recently completed or running jobs

All slurm commands have additional options available in their manual pages

Type `"man <command>"` to bring up the man page

Type `"q"` to exit the man page

# M | BIOSTATISTICS

OnDemand provides an integrated, single access point for all of your HPC resources.

[Home](#) / [My Interactive Sessions](#) / [RStudio](#)

### Interactive Apps

[Overview](#)
[Interactive App](#)
[Sessions](#)
[Jobs](#)
[Quotas](#)
[Interactive App \(Default\)](#)
[Sessions](#)
[Jobs](#)
[Quotas](#)

## RStudio

This app will launch RStudio Server on the RStudio OnDemand cluster.

RStudio

This app will launch RStudio Server on the RStudio OnDemand cluster.

Number of hours

1

Memory

0

Memory (GB)

Number of cores

1

☐ Load last saved workspace (.RData) when session starts

☐ I would like to receive an email when the session starts

Launch

\* All RStudio session data is generated and stored under the user's home directory in the corresponding user's file system.



# Basic Linux Commands

ls - list files or directories

cd - change directory

pwd - print working directory

cat - concatenate

cp - copy

mv - move

rm - remove

mkdir - make a directory

# Submit Jobs

Should always submit a job to use the cluster. Do not do too much computation on the login nodes. Your session will be killed off.

If you need to do interactive work, submit an interactive job

```
srun --pty /bin/bash
```