

# Rbootcamp

Michael Kleinsasser

8/8/2019

# Mike's Personal Introduction

- R programmer for the Department of Biostatistics
- Write and maintain R packages for faculty and students
- Consult faculty and students on writing R packages, optimization



# Rbootcamp Introduction

## Goals:

- R basics: syntax, common functions, etc., via Rstudio
- R functions and for loops (functions and advanced control structures)
- Basics of R on the cluster (non-interactive R using BATCH scripts)
- Simulation analysis with comment on efficiency

# Materials

- All bootcamp materials online at <https://github.com/umich-biostatistics/Rbootcamp>
  - Handouts for each topic with examples to work through
  - R scripts of our examples
- Go to link and download zip archive, extract

# Setup

Go to Rstudio cloud to follow along:

- Enter username, etc. for free account
- Follow along by typing commands in my slides
- If you have Rstudio/R, open that
- Recommended: Install R/Rstudio for days 2, 3
- See course materials for download instructions

# R Basics: Big Picture

- R is a sophisticated calculator for statistics

Chambers (2016) Extending R:

- Everything that exists in R is an object
- Everything that happens in R is a function call

Obtain a basic working knowledge of R objects and functions,

- Google the rest

# R Basics: a basic schematic view

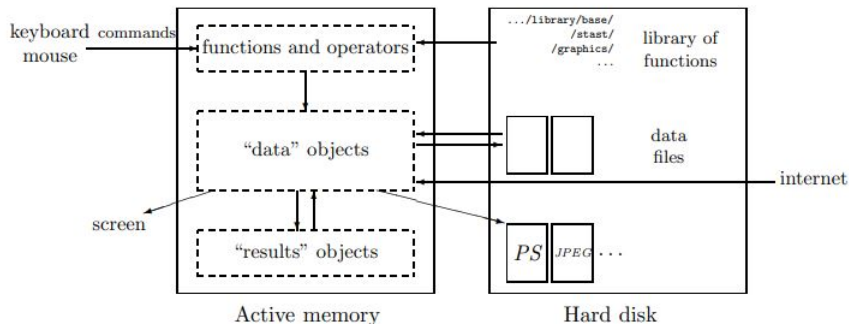


Figure 1: A schematic view of how R works.

# R Basics: Goals

- Data types and functions
  - Create object having data types
  - Combine those into data structures
  - Write basic R function
- Learn parts of R most useful to statisticians
  - How do most modeling functions work in R, and
  - How to inspect structure and content of objects
- Apply knowledge to simulation
  - How do most modeling functions work in R, and
  - How to inspect structure and content of objects



# Use R as a calculator

Standard operations:

```
# multiply *, divide /, add + subtract -  
18056.983 - 1005.118 + 22.53
```

```
## [1] 17074.4
```

```
( (pi - 3.14) / (3.14) ) * 100
```

```
## [1] 0.05072145
```



# Operators: arithmetic and logical

Operator	Description
+	addition
-	subtraction
/	division
^	exponential
%%	modulus (x mod y)
%/%	integer division

**Table 1:** Arithmetic Operators

Operator	Description
<	less than
<=	less than or equal
>	greater than
>=	greater than or equal
==	exactly equal to
!=	not equal to
!x	not x (x logical)
x y	x OR y
x&y	x AND y
isTRUE(x)	is x TRUE

**Table 2:** Logical Operators

# Operators: an example

Test if these two expressions are equivalent in R:

# Operators: an example

Test if these two expressions are equivalent in R:

Convert the following expression to R code:

# Create, list, delete objects

Create object with “assign” operator

- arrow then minus sign <-
- single equal sign =

```
# x gets the number 3.14
```

```
x <- 3.14
```

```
x      # print x
```

```
## [1] 3.14
```

```
# equivalently
```

```
x = 3.14
```

```
x      # print x
```

```
## [1] 3.14
```

# Create, list, delete objects

Objects we create are stored in memory, e.g.:

```
name = "Carmen"  
n1 = 10  
n2 = 100  
m = 0.5
```

Use `ls()` function to list all objects in memory:

```
ls()
```

```
## [1] "m"      "n1"     "n2"     "name"   "x"
```

Notice: I created `x` before, it's still in memory.

# Create, list, delete objects

Use `ls()` function to list all objects in memory:

```
ls()
```

```
## [1] "m"      "n1"     "n2"     "name"   "x"
```

The function `ls.str()` displays some details about objects in memory:

```
ls.str()
```

```
## m :   num 0.5
## n1 :   num 10
## n2 :   num 100
## name :  chr "Carmen"
## x :   num 3.14
```

# Create, list, delete objects

To delete objects in memory, use `rm()` function

```
rm(x) # deletes object named x  
ls()  # which objects remain in memory?
```

```
## [1] "m"      "n1"      "n2"      "name"
```

```
rm(m, n2, name) # remove multiple objects  
ls()
```

```
## [1] "n1"
```

```
rm(list = ls()) # remove everything from memory  
ls()
```

```
## character(0)
```



# The on-line help

R has structured help pages providing “how-to”

- **Description:** what function does
- **Usage:** name with arguments and options
- **Arguments:** how each argument should be structured
- **Details:** more detailed description
- **Value:** How the output is structured/ what it contains
- **Examples:** examples of the function in use

```
?rm           # help documentation for rm function  
help("rm")    # alternately
```

remove {base}

R Documentation

## Remove Objects from a Specified Environment

### Description

`remove` and `rm` can be used to remove objects. These can be specified successively as character strings, or in the character vector `list`, or through a combination of both. All objects thus specified will be removed.

If `envir` is `NULL` then the currently active environment is searched first.

If `inherits` is `TRUE` then parents of the supplied directory are searched until a variable with the given name is encountered. A warning is printed for each variable that is not found.

### Usage

```
remove(..., list = character(), pos = -1,  
        envir = as.environment(pos), inherits = FALSE)  
  
rm      (... , list = character(), pos = -1,  
        envir = as.environment(pos), inherits = FALSE)
```

## Other R help

Many package writers create Vignettes and READMEs

How to view vignettes?

```
vignette(all = TRUE) # list vignettes for installed packages  
vignette(all = FALSE) # vignettes from attached packages
```

How to view READMEs?

Other help...

# Objects in R

Objects (data, model output, functions)

- Characterized by their **names** and **content**
- attributes - specify the kind of data represented
  - e.g. mode, length

```
x = 1  
mode(x)
```

```
## [1] "numeric"
```

```
length(x)
```

```
## [1] 1
```

# Basic data “modes” in R

The mode is the basic type of the elements of an object

The four main modes:

- numeric, comes in two flavors: integer, numeric
- character
- logical
- complex

```
num = 15.533; name = "Mike"; isStudent = TRUE
```

```
mode(num); mode(name); mode(isStudent)
```

```
## [1] "numeric"
```

```
## [1] "character"
```

```
## [1] "logical"
```

# Atomic vectors

Fundamental data structure in R:

- atomic vector - vector in which every element is of same mode

To create an atomic vector, use `c()` function:

```
c(3.145, 2.18, 9.98e3, 0.05)
```

```
## [1]      3.145      2.180 9980.000      0.050
```

Example: create empty character vector of length 3 and store your full name

```
vector(mode = "character", length = 3)
```

```
## [1] "" "" ""
```

# Data types examples

3 ways to create numeric vector:

```
# empty numeric vector  
y1 <- numeric(6)  
y1      # print y1
```

```
## [1] 0 0 0 0 0 0
```

```
y2 <- vector(mode = "numeric", length = 6)  
y2      # print y2
```

```
## [1] 0 0 0 0 0 0
```

```
y3 <- c(5, 13.222, 2, 0.001, 77.4, 31.9)  
y3      # print y3
```

```
## [1] 5.000 13.222 2.000 0.001 77.400 31.900
```

# Objects in R: NA

NA means “Not Available” and it denotes missing data  
(Insert after atomic vectors introduced)

```
c(3, 5, 9, NA, 18, 25, NA)
```

```
## [1] 3 5 9 NA 18 25 NA
```



# R data types/structures

Four fundamental data types:

- character, numeric (numeric or integer), logical, complex

Combine to form data structures

- atomic vector (atomic - vector of single type)
- list
- matrix
- data.frame
- factor

We will focus on matrices and data.frames

# Matrices

Matrices are the natural extension of atomic vectors into 2 dimensions

- any mode can be used, but numeric most common:

Syntax:

```
m = matrix(data, nrow, ncol, byrow, dimnames)
```

- **data** is the input vector which becomes the data elements of the matrix
- **nrow, ncol** is the number of rows/columns to be created
- **byrow** is a logical clue. If TRUE then the input vector elements are arranged by row
- **dimnames** is the names assigned to the rows and columns

# Matrix examples:

Identity matrix:

```
dat = c(1,0,0,0,1,0,0,0,1)  # data
iden = matrix(data = dat, nrow = 3, byrow = T)
iden  # print matrix
```

```
##      [,1] [,2] [,3]
## [1,]    1    0    0
## [2,]    0    1    0
## [3,]    0    0    1
```

Easier:

```
iden = diag(rep(1,3))
iden
```

## access elements of a matrix

- single brackets used to access elements

Access individual elements:

```
# 2x5 matrix of numbers 1 to 10  
P = matrix(data = 1:10, nrow = 2)  
P[1,3] # row 1, column 3
```

```
## [1] 5
```

```
P[nrow(P),ncol(P)] # row 2, column 5 (bottom right position)
```

```
## [1] 10
```

Access entire rows/columns:

```
P[,3];
```

```
## [1] 5 6
```

# The data.frame

The **data.frame** is the most common way to store and work with data in R

- Not surprising: they are designed for this purpose
- Most modeling functions work on data.frames

Composed of a list of equal length atomic vectors (can be of any type)

Example:

The following are data on students in the class:

- Has Master's (logical): TRUE FALSE FALSE TRUE
- GPA (numeric): 3.1 4.0 2.9 3.6
- First Name (character): Mike Dan Sara Karen

Convert to three atomic vectors of appropriate type.

# data.frame examples

Create a data.frame out of the following “class” data:

- Has Master's (logical): TRUE FALSE FALSE TRUE
- GPA (numeric): 3.1 4.0 2.9 3.6
- First Name (character): Mike Dan Sara Karen

```
# store data
has_ms <- c(TRUE, FALSE, FALSE, TRUE)
gpa <- c(3.1, 4.0, 2.9, 3.6)
name <- c("Mike", "Dan", "Sara", "Karen")
# Create data.frame
dat <- data.frame(has_MS = has_ms, GPA = gpa, Name = name)
dat      # print data.frame
```

```
##   has_MS GPA  Name
## 1   TRUE 3.1  Mike
## 2  FALSE 4.0   Dan
## 3  FALSE 2.9   Sara
## 4   TRUE 3.6 Karen
```

## access elements of a data.frame

Each vector of a data.frame contains the values of a variable

Access each vector with the dollar sign \$

Ex: Extract the GPA column and print it

```
dat$GPA
```

```
## [1] 3.1 4.0 2.9 3.6
```

Another example: ToothGrowth data.

```
ToothGrowth$len
```

```
## [1] 4.2 11.5 7.3 5.8 6.4 10.0 11.2 11.2 5.2 7.0 16.5
## [15] 22.5 17.3 13.6 14.5 18.8 15.5 23.6 18.5 33.9 25.5 26.4
## [29] 23.3 29.5 15.2 21.5 17.6 9.7 14.5 10.0 8.2 9.4 16.5
## [43] 23.6 26.4 20.0 25.2 25.8 21.2 14.5 27.3 25.5 26.4 22.4
## [57] 26.4 27.3 29.4 23.0
```

# data.frame

Preview head (first few rows) of data.frame:

```
head(ToothGrowth)
```

```
##      len supp dose
## 1  4.2   VC  0.5
## 2 11.5   VC  0.5
## 3  7.3   VC  0.5
## 4  5.8   VC  0.5
## 5  6.4   VC  0.5
## 6 10.0   VC  0.5
```

View tail of data.frame:

```
tail(dat)
```

View entire data.frame in new window:

```
View(ToothGrowth)
```



# Inspect an object

- `class()` - what kind of object is it (high-level)?
- `typeof()` - what is the data type (low-level)?
- `length()` - how long is it?
- `attributes()` - does it have meta-data?

# Inspect an object

- `class()` - what kind of object is it (high-level)?

```
class(ToothGrowth)
```

```
## [1] "data.frame"
```

- `typeof()` - what is the data type (low-level)?

```
typeof(ToothGrowth$supp)
```

```
## [1] "integer"
```

# Inspect an object

- `length()` - how long is it?

```
length(ToothGrowth$dose)
```

```
## [1] 60
```

- `attributes()` - does it have meta-data?

```
attributes(ToothGrowth)
```

```
## $names
```

```
## [1] "len" "supp" "dose"
```

```
##
```

```
## $class
```

```
## [1] "data.frame"
```

```
##
```

```
## $row.names
```

```
## [1] 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18
```

# R functions

R function syntax:

```
NAME <- function(ARG1, ARG2, ARG3) {  
  DO SOMETHING  
  STORE RESULT  
  return(RESULT)  
}
```

```
pow <- function(base, expon) { # power function  
  prod(rep(base, expon)) # base^(expon)  
}  
# Use power function  
pow(5, 2)
```

```
## [1] 25
```

```
pow(10, 3)
```

```
## [1] 1000
```

# Common R functions

R has a huge collection of packages:

- 6,000+ packages for data analysis build (on CRAN alone)

Example: `lm` (linear models)

- Use `?lm` to read help documentation

`lm {stats}`

R Documentation

## Fitting Linear Models

### Description

`lm` is used to fit linear models. It can be used to carry out regression, single stratum analysis of variance and analysis of covariance (although [aov](#) may provide a more convenient interface for these).

### Usage

```
lm(formula, data, subset, weights, na.action,  
   method = "qr", model = TRUE, x = FALSE, y = FALSE, qr = TRUE,  
   singular.ok = TRUE, contrasts = NULL, offset, ...)
```

### Arguments

- |                      |  |
|----------------------|--|
| <code>formula</code> | an object of class " <a href="#">formula</a> " (or one that can be coerced to that class): a symbolic description of the model to be fitted. The details of model specification are given under 'Details'.   |
| <code>data</code>    | an optional data frame, list or environment (or object coercible by <a href="#">as.data.frame</a> to a data frame) containing the variables in the model. If not found in <code>data</code> , the variables are taken from <code>environment(formula)</code> , typically the environment from which <code>lm</code> is called. |

# Fit a linear model with lm

- Use built-in data set ToothGrowth
- ?ToothGrowth for help:

## The Effect of Vitamin C on Tooth Growth in Guinea Pigs

### Description

The response is the length of odontoblasts (cells responsible for tooth growth) in 60 guinea pigs. Each animal received one of three dose levels of vitamin C (0.5, 1, and 2 mg/day) by one of two delivery methods, orange juice or ascorbic acid (a form of vitamin C and coded as VC).

### Usage

```
ToothGrowth
```

### Format

A data frame with 60 observations on 3 variables.

```
[,1] len    numeric Tooth length  
[,2] supp   factor   Supplement type (VC or OJ).  
[,3] dose   numeric Dose in milligrams/day
```

# View the data

View data in new window:

```
View(ToothGrowth)
```

Or use head to view only first 6 rows:

```
head(ToothGrowth)
```

```
##      len supp dose
## 1  4.2   VC  0.5
## 2 11.5   VC  0.5
## 3  7.3   VC  0.5
## 4  5.8   VC  0.5
## 5  6.4   VC  0.5
## 6 10.0   VC  0.5
```

How big is the data?

```
dim(ToothGrowth)
```

```
## [1] 60  3
```

# Using lm() function for linear models

Call lm on the data and formula, store result "lm" object:

```
tooth_fit = lm(formula = len ~ supp + dose,  
               data = ToothGrowth)
```

Formulas in R:

```
len ~           # Response column name, ~ for "="  
  supp +        # First predictor name + for "+"  
  dose          # second predictor name
```

Many R functions use the formula argument.



# Getting detailed information

Basic “print” of model:

```
print(tooth_fit)      # equivalent to tooth_fit

##
## Call:
## lm(formula = len ~ supp + dose, data = ToothGrowth)
##
## Coefficients:
## (Intercept)      suppVC          dose
##      9.272      -3.700      9.764
```

Detailed summary:

```
summary(tooth_fit)

##
## Call:
## lm(formula = len ~ supp + dose, data = ToothGrowth)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -6.600 -3.700  0.373  2.116  8.800
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   9.2725     1.2824   7.231 1.31e-09 ***
## suppVC       -3.7000     1.0936  -3.383  0.0013 **
## dose         9.7636     0.8768  11.135 6.31e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
```

# Understanding R classes

- What is this thing?

```
class(tooth_fit)
```

- What are the methods for this object?

```
methods(class = "lm")
```

- What is its structure? (i.e., what's in it)

```
str(tooth_fit)
```

# Understanding R classes

- What is this thing?

```
class(tooth_fit)
```

```
## [1] "lm"
```

# Understanding R classes

- What are the methods for this object?

```
methods(class = "lm")
```

```
## [1] add1          alias          anova          case.names
## [5] coerce        confint        cooks.distance deviance
## [9] dfbeta        dfbetas        drop1          dummy.coef
## [13] effects       extractAIC     family         formula
## [17] hatvalues     influence      initialize     kappa
## [21] labels        logLik         model.frame    model.matrix
## [25] nobs          plot           predict        print
## [29] proj          qr             residuals      rstandard
## [33] rstudent      show           simulate       slotsFromS3
## [37] summary       variable.names vcov
## see '?methods' for accessing help and source code
```

# Understanding R classes

- What is its structure? (i.e., what's in it)

```
str(tooth_fit)
```

```
## List of 13
## $ coefficients : Named num [1:3] 9.27 -3.7 9.76
## ..- attr(*, "names")= chr [1:3] "(Intercept)" "suppVC" "dose"
## $ residuals    : Named num [1:60] -6.25 1.05 -3.15 -4.65 -4.05 ...
## ..- attr(*, "names")= chr [1:60] "1" "2" "3" "4" ...
## $ effects      : Named num [1:60] -145.73 14.33 47.16 -3.86 -3.26 ...
## ..- attr(*, "names")= chr [1:60] "(Intercept)" "suppVC" "dose" "" ...
## $ rank         : int 3
## $ fitted.values: Named num [1:60] 10.5 10.5 10.5 10.5 10.5 ...
## ..- attr(*, "names")= chr [1:60] "1" "2" "3" "4" ...
## $ assign       : int [1:3] 0 1 2
## $ qr          :List of 5
## ..$ qr       : num [1:60, 1:3] -7.746 0.129 0.129 0.129 0.129 ...
## .. ..- attr(*, "dimnames")=List of 2
## .. .. ..$ : chr [1:60] "1" "2" "3" "4" ...
## .. .. ..$ : chr [1:3] "(Intercept)" "suppVC" "dose"
## .. ..- attr(*, "assign")= int [1:3] 0 1 2
## .. ..- attr(*, "contrasts")=List of 1
## .. .. ..$ supp: chr "contr.treatment"
## ..$ qraux: num [1:3] 1.13 1.11 1.11
## ..$ pivot: int [1:3] 1 2 3
## ..$ tol : num 1e-07
## ..$ rank : int 3
## ..- attr(*, "class")= chr "qr"
## $ df.residual : int 57
## $ contrasts    :List of 1
## ..$ supp: chr "contr.treatment"
## $ xlevels     :List of 1
```

# Understanding R classes

- Pull something out of the “lm” fit object:

```
tooth_fit$fitted.values      # y_hat's for the linear model
```

```
##      1      2      3      4      5      6      7      8
## 10.45429 10.45429 10.45429 10.45429 10.45429 10.45429 10.45429 10.45429
##      9     10     11     12     13     14     15     16
## 10.45429 10.45429 15.33607 15.33607 15.33607 15.33607 15.33607 15.33607
##     17     18     19     20     21     22     23     24
## 15.33607 15.33607 15.33607 15.33607 25.09964 25.09964 25.09964 25.09964
##     25     26     27     28     29     30     31     32
## 25.09964 25.09964 25.09964 25.09964 25.09964 25.09964 14.15429 14.15429
##     33     34     35     36     37     38     39     40
## 14.15429 14.15429 14.15429 14.15429 14.15429 14.15429 14.15429 14.15429
##     41     42     43     44     45     46     47     48
## 19.03607 19.03607 19.03607 19.03607 19.03607 19.03607 19.03607 19.03607
##     49     50     51     52     53     54     55     56
## 19.03607 19.03607 28.79964 28.79964 28.79964 28.79964 28.79964 28.79964
##     57     58     59     60
## 28.79964 28.79964 28.79964 28.79964
```

# Extracting data from model objects

Some generic extraction methods:

```
coef(tooth_fit)           # model coefficients
```

```
coef(summary(tooth_fit))  # adds test statistics, p-values
```

```
vcov(tooth_fit)           # variance/covariance matrix
```

Note: depending on implementation, these may not be available - Check methods with “methods(object)” before attempting

# Extracting data from model objects (in detail)

Extract coefficients, test stats, and p-values

```
coef(summary(tooth_fit))    # adds test statistics, p-values
```

	Estimate	Std. Error	t value	Pr(> t )
## (Intercept)	9.272500	1.2823649	7.230781	1.312335e-09
## suppVC	-3.700000	1.0936045	-3.383307	1.300662e-03
## dose	9.763571	0.8768343	11.135025	6.313519e-16



# Predict new values

- `predict()` function is generic and works with many models
- Pass in a new `data.frame` with the same column names:

```
to_predict = data.frame(dose = 0.5, supp = "VC")
```

```
predict(tooth_fit, newdata = to_predict)
```

```
##           1  
## 10.45429
```

## Predict new values (example 2)

- `predict()` function is generic and works with many models
- Pass in a new `data.frame` with the same column names:

```
to_predict = data.frame(dose = seq(0,1,0.1), supp = "OJ")
```

```
predict(tooth_fit, newdata = to_predict)
```

```
##           1           2           3           4           5           6           7
##  9.27250 10.24886 11.22521 12.20157 13.17793 14.15429 15.13064 16.10708
##           9          10          11
## 17.08336 18.05971 19.03607
```


# R on cluster (non-interactive R)

## Cluster Computation

- Dan Barker danbarke at umich.edu
- Cluster System Administrator

# R workflow for statistical analysis

- Bootcamps throughout semester on each area



# Advanced control structures

Common misconception: “loops in R are slow”

Explain

## Loops: often not necessary

- Many R functions are “vectorized” (vector in, vector out)
- While most other languages require loops, R does not:

```
a = c(5, 2, 4, 12, 1)
```

```
b = c(2, 0, 3, -1, 2)
```

```
a + b      # vector + vector = vector
```

```
## [1]  7  2  7 11  3
```

Takehome: When possible, operate on vectors and matrices, don't loop over each row/position index

## Problem: not all functions are vectorized

- Some functions, like `read.table()` for reading a table of data into R, are not vectorized
- But what if we have a list of files to read in? “data1.txt”, “data2.txt”, “data3.txt”, ..., “data50.txt”

```
# the 50 data set names  
file_names = paste0("data", 1:50, ".txt")
```

Attempt it, will cause error:

```
read.table(file_names) # error!
```

## map() from purrr package to avoid loop

- map() allows you to apply a function to each element of a vector
- faster, easier to read than a loop

```
read.table(file_names)  # error!
```

```
  # list of 50 data sets
my_dat_list = map(file_names, read.table)
  # results from reading data1.txt
my_dat_list[[1]]
  # results from reading data50.txt
my_dat_list[[50]]
```



# Create your own functions for map

Generate random samples from a Normal with different variances

```
draws = map(2:20, function(x) rnorm(25, mean=0, sd=x))  
# returns a list with vectors of draws
```

```
# first vector of draws:
```

```
draws[[1]]
```

```
# last vector of draws:
```

```
draws[[19]]
```

# Create your own functions for map

Now, estimate the standard error:

```
map(draws, sd)  # sd() is standard deviation in R  
map_dbl(draws, sd) # numeric vector
```

# For loop vs map() version

- Standard for loop:
- map():

# Sometimes loops are required

Sometime you just need loops:

- Growth model, new values depend on previous values

```
N = 20
for (i in 2:30) {
  f = rpois(1, 0.15*N[i-1])    # births
  d = rbinom(1, N[i-1], 0.1)   # deaths
  N[i] = N[i-1] + f - d
}

plot(seq_along(N), N)
```

Show the plot!

## Problem: Growing object is slow

```
N = 20
for (i in 2:30) {
  f = rpois(1, 0.15*N[i-1])    # births
  d = rbinom(1, N[i-1], 0.1)  # deaths
  N[i] = N[i-1] + f - d
}

plot(seq_along(N), N)
```

- Our growth model loop is slow because we are growing a vector at each iteration
- Solution pre-allocate vector (or any data type), then fill with loop

# Efficient memory usage

Improved code:

```
# Pre-allocate to correct size
N = vector(mode = "numeric", length = 30)
N[1] = 20
for (i in 2:30) {
  f = rpois(1, 0.15*N[i-1]) # births
  d = rbinom(1, N[i-1], 0.1) # deaths
  N[i] = N[i-1] + f - d
}

plot(seq_along(N), N)
```

# overview of loops

- Although slow at times, writing code is a trade-off between time consumed by the programmer and time gained with efficient code
- Most times, loops are fast (enough) and convenient
- Strategy: Write a loop, then if slow, re-write
- Use vectorized functions whenever possible
- Map when function is not vectorized
  - Avoids the need to track indexes
  - Does the pre-allocation for you
- If you use loops, avoid growing lists/collections/data.frames
  - pre-allocate memory

# Simulation

Many built-in distribution functions in R:

They have prefix r, q, p, ... for random, ...

- probability density: d\_\_\_\_\_
- cumulative probability density p\_\_\_\_\_
- quantile: q\_\_\_\_\_
- random draw: r\_\_\_\_\_

Example: uniform unif



# Simulation

distribution	function
Normal	<code>rnorm(n, mean=0, sd=1)</code>
exponential	<code>rexp(n, rate=1)</code>
uniform	<code>runif(n, min=0, max=1)</code>
gamma	<code>rgamma(n, shape, scale=1)</code>
poisson	<code>rpois(n, lambda)</code>
Weibull	<code>rweibull(n, shape, scale=1)</code>
Cauchy	<code>rcauchy(n, location=0, scale=1)</code>
beta	<code>rbeta(n, shape1, shape2)</code>
Student t	<code>rt(n, df)</code>
binomial	<code>rbinom(n, size, prob)</code>
logistic	<code>rlogis(n, location=0, scale=1)</code>


**Table 3:** Built-in distributions

# Simulation: Confidence intervals

Open simulations.Rmd Rmarkdown document in course files

Exercise: Is the coverage rate of confidence intervals for the mean accurate?

- Given conditions similar to ours, does our method for generating confidence intervals capture the true value 95% of the time?



# Simulate Im

Need to come up with some question to answer.

# Bootstrap lm results

# Useful R packages

- ggplot2 - Best package for plotting data
  - getting started:
- dplyr - Best package for manipulating/processing data
  - getting started:

# Future bootcamps

- dplyr: date \_\_\_\_\_
- ggplot2: date \_\_\_\_\_